

Lab 2 – GPT from scratch

IASD / MASH – LLMs course

Florian Le Bronnec

November 12, 2024

Table of Contents

① Lab 2 – Recap

Last time

Questions / remarks

② Training

③ Decoding

Table of Contents

① Lab 2 – Recap

Last time

Questions / remarks

② Training

③ Decoding

Last time

Goal: Implement a version of the GPT model from scratch.

- ① We built the **attention module**.
- ② We built the **transformer block**.
 - Attention
 - Feedforward
 - LayerNorm + skip connection
- ③ We built the **transformer model** by stacking transformer blocks.
- ④ We added a linear layer to the output of the transformer model for the **vocabulary projection**.

Layer norm

- **Error in the code:** layer norm shouldn't use detach gradients!

Layer norm

- **Error in the code:** layer norm shouldn't use detach gradients!
- Where should we put the layer norm? In recent LLaMA models, **pre-layer normalization** is applied.

$$\mathbf{y} = \text{SelfAttention}(\text{LayerNorm}(\mathbf{x})) + \mathbf{x}$$

$$\mathbf{z} = \text{FeedForward}(\text{LayerNorm}(\mathbf{y})) + \mathbf{y}$$

In contrast, the original Transformer model uses **post-layer normalization**:

$$\mathbf{y} = \text{LayerNorm}(\text{SelfAttention}(\mathbf{x}) + \mathbf{x})$$

- **Rule of thumb:** Apply layer normalization before each sublayer (like attention or feed-forward) in deep or autoregressive models, such as GPT-style transformers.

Attention's scaling

- **Add scaling:** $\frac{1}{\sqrt{d}}$ (forgotten in the previous lab).

Masks in multi-head self-attention

Goal: Ensure that the model only attends to relevant tokens by:

- Preventing access to future tokens (important for generation tasks)
- Ignoring padding tokens that don't contain information

Causal masking (masking future tokens)

- **Why?** In text generation, each position should only access past tokens to make sure predictions are causal.
- **How?** We use an upper triangular matrix to mask future tokens:

$$\text{mask}_{\text{causal}} = \text{triu}(\mathbf{1}_{\text{shape}(\mathbf{s}_{ij})}, 1)$$

- **Setting scores to $-\infty$:** Scores for future tokens are set to $-\infty$. Then a_{ij} will be zero for these tokens.

Masks in multi-head self-attention

Padding masking (ignoring padding tokens)

- **Why?** Padding tokens are added to make all inputs the same length, but they don't contain any information. The model should ignore them.
- **How?** The padding mask is calculated using `attention_mask`.
- **Setting scores to $-\infty$:** Similar to causal masking, padding tokens are set to $-\infty$ in \mathbf{s}_{ij} .

Combining causal and padding masks

- **Why combine both?** Combining both masks ensures that each token attends only to valid, relevant tokens, ignoring both future and padding tokens.
- **Final adjustment:** We add both masks to \mathbf{s}_{ij} and scale them by $-\infty$:

$$\mathbf{s}_{ij} = \mathbf{s}_{ij} + (\text{mask}_{\text{causal}} + \text{mask}_{\text{padding}}) \times (-\infty)$$

$$\mathbf{a}_{ij} = \text{softmax}(\mathbf{s}_{ij}, \text{dim} = -1)$$

Making equal size tensors (padding)

- **Error in the code:** Typo when using a wrong variable name (use `label_ids`)

Table of Contents

① Lab 2 – Recap

Last time

Questions / remarks

② Training

③ Decoding

Step 1: Input Representation as Word IDs

- The input sequence is represented by a list of unique IDs for each token, starting with a special BOS (beginning-of-sequence) token.

Token	Word ID
<BOS>	0
Paris	102
is	76
the	34
capital	159
of	45
France	203

Step 2: Mapping Word IDs to Embedding Vectors

- Each token ID is converted into a vector in the embedding space.
- Example embeddings are shown as 3-dimensional vectors.

Token	Word ID	Embedding Vector
<BOS>	0	[0.1, 0.0, 0.3]
Paris	102	[0.2, 0.1, 0.4]
is	76	[0.3, 0.6, 0.5]
the	34	[0.5, 0.2, 0.7]
capital	159	[0.6, 0.4, 0.3]
of	45	[0.1, 0.8, 0.2]
France	203	[0.4, 0.9, 0.1]

Step 3: Causal Processing with Self-Attention

- Causal processing is achieved through masked self-attention.
- Each token can only attend to itself and previous tokens.

Self-Attention with Mask

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} + M \right) V$$

where Q , K , and V are the query, key, and value matrices, d_k is the dimension of keys, and M is the causal mask.

- The causal mask M is set to $-\infty$ for future tokens, blocking attention to future positions.

Step 4: Logits and Prediction for Next Token

- After processing, each token produces a logits vector representing the probability distribution for the next word.
- Example logits for each position (showing top 3 predictions):

Token	Top Logits for Next Token
<BOS>	{Paris: 0.7, London: 0.2, Berlin: 0.1}
Paris	{is: 0.6, was: 0.2, and: 0.1}
is	{the: 0.5, a: 0.3, an: 0.1}
the	{capital: 0.4, city: 0.3, center: 0.2}
capital	{of: 0.7, in: 0.2, for: 0.1}
of	{France: 0.6, Europe: 0.3, Paris: 0.1}
France	{<END>: 0.9, is: 0.05, the: 0.02}

Step 5: Teacher Forcing and Loss Calculation

- In teacher forcing, we calculate the loss as the negative log probability of the true next token.

Token	Should predict	Logits (Top Scores)	Loss
<BOS>	Paris	{Paris: 0.7, London: 0.2}	$-\log(0.7)$
Paris	is	{is: 0.6, was: 0.2}	$-\log(0.6)$
is	the	{the: 0.5, a: 0.3}	$-\log(0.5)$
the	capital	{capital: 0.4, city: 0.3}	$-\log(0.4)$
capital	of	{of: 0.7, in: 0.2}	$-\log(0.7)$
of	France	{France: 0.6, Europe: 0.3}	$-\log(0.6)$
France	<END>	{<END>: 0.9, is: 0.05}	$-\log(0.9)$

Summary: From Input to Teacher Forcing in LLMs

- Convert tokens to IDs with BOS.
- Map IDs to embeddings.
- Process embeddings causally through masked self-attention.
- Generate logits for each position.
- Apply teacher forcing, compute log probability, and backpropagate.

Self-Attention Recap

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} + M \right) V$$

where M is the causal mask to restrict attention to preceding tokens only.

Table of Contents

① Lab 2 – Recap

Last time

Questions / remarks

② Training

③ Decoding

Introduction to Inference

- Inference is the process by which a model generates tokens to extend an input sequence.
- This is done iteratively, selecting one token at a time and appending it to the sequence.
- Two common methods are:
 - Greedy Decoding
 - Ancestral Sampling

Inference in classification

At train time

Classification models are trained with the MLE objective, i.e., they maximize $P_{\theta}(y = \text{class} \mid x)$.

At inference time

For an input x , the model gives a probability distribution over the classes $P_{\theta}(\cdot \mid x)$.

Then you fix a decision rule, usually:

$$\hat{y} = \arg \max_y P_{\theta}(y \mid x). \quad (1)$$

\implies For classification models, i.e., **encoders** in NLP, we do the same.

Inference in generation

At train time

For generative models, we maximize instead the factorized density:
 $\prod_{i=1}^L P_{\theta}(y_i = w_i \mid w_{<i}).$

At inference time

Is the following decision rule still a good choice?

$$\hat{y} = \arg \max_y P_{\theta}(y \mid x) = \arg \max_{y_1, \dots, y_L} \prod_{i=1}^L P_{\theta}(y_i \mid y_{<i}).$$

No! We are taking the arg max over V^L combinations, highly **intractable**.

Inference in generation

Workaround: approximate $\arg \max_{y_1, \dots, y_L} \prod_{i=1}^L P_{\theta}(y_i \mid y_{<i})$ with a **greedy algorithm**.

Simply:

- $\hat{y}_1 = \arg \max_{y_1} P_{\theta}(y_1),$
- $\hat{y}_2 = \arg \max_{y_2} P_{\theta}(y_2 \mid \hat{y}_1),$
- \vdots
- $\hat{y}_i = \arg \max_{y_i} P_{\theta}(y_i \mid \hat{y}_{<i}).$

\implies At each step, the $\arg \max$ is only performed over V possibilities.

Greedy Decoding

Definition

Greedy decoding selects the token with the highest probability at each step, creating a deterministic sequence.

- **Step 1:** Start with an input sequence, e.g., "Paris is the".
- **Step 2:** The model predicts probabilities for the next token:
$$\{\text{capital} : 0.6, \text{city} : 0.3, \text{center} : 0.1\}$$
- **Step 3:** Select the token with the highest probability ("capital").
- **Step 4:** Append the selected token to the sequence:

"Paris is the capital"

- **Step 5:** Repeat Steps 2-4 until reaching the end token (<END>) or desired length.

Greedy Decoding - Iterative Growth Example

- Initial sequence: "Paris is the"

- **Iteration 1:** Predict "capital"

"Paris is the capital"

- **Iteration 2:** Predict "of"

"Paris is the capital of"

- **Iteration 3:** Predict "France"

"Paris is the capital of France"

- Sequence reaches an end token (<END>), completing generation.

Ancestral Sampling

Definition

Ancestral sampling selects tokens based on their probabilities, allowing randomness and creating diverse possible outputs.

- **Step 1:** Start with an input sequence, e.g., "Paris is the".
- **Step 2:** The model predicts probabilities for the next token:
$$\{\text{capital} : 0.6, \text{city} : 0.3, \text{center} : 0.1\}$$
- **Step 3:** Sample a token based on these probabilities, e.g., "city" (with 30% chance).
- **Step 4:** Append the sampled token to the sequence:

"Paris is the city"

- **Step 5:** Repeat Steps 2-4 until reaching <END> or desired length.

Ancestral Sampling - Iterative Growth Example

- Initial sequence: "Paris is the"
- **Iteration 1:** Sample "city"

"Paris is the city"

- **Iteration 2:** Sample "of"

"Paris is the city of"

- **Iteration 3:** Sample "dreams"

"Paris is the city of dreams"

- Ancestral sampling introduces variety, so the final output could vary on each run.

Comparison of Greedy Decoding and Ancestral Sampling

- **Greedy Decoding:**

- Deterministic, choosing only the highest-probability token at each step.
- Suitable for generating the most probable and consistent sequence.

- **Ancestral Sampling:**

- Stochastic, selecting tokens based on probabilities, allowing randomness.
- Produces varied and potentially more creative sequences.

Summary

- Inference in LLMs can be done using methods like greedy decoding and ancestral sampling.
- **Greedy decoding** is deterministic and produces an approximated most likely sequence.
- **Beam search** is a generalization of greedy decoding that considers multiple likely sequences.
- **Ancestral sampling** is stochastic.