

Structure, Attention, and BERT

Alexandre Allauzen

Fall 2024



Roadmap

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Outline

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

The simplest classifier based on word embeddings

Input text is now a sequence of vectors (embeddings)

$$\mathbf{X} \text{ or } \mathbf{X}^t = \underbrace{\begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline \end{array}}^d \left. \vphantom{\begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline \end{array}} \right\} L \text{ words}$$

Derive a vector of features that represent the text

$$\mathbf{h} = \sum_{i=1}^L \underbrace{\mathbf{x}_i}_{\text{emb. of word } i}$$

Classification

$\text{softmax}(\mathbf{W}^o \mathbf{h})$ (multiclass) or $\sigma(\mathbf{w}^o \mathbf{h})$ (binary)

Limitations of BOW classifier

$$\mathbf{h} = \sum_{i=1}^L \underbrace{\mathbf{x}_i}_{\text{emb. of word } i}$$

Limitations

- Words are equally important
- Word order independent
- Miss contextual information (local/global)

Local contexts

the	end	is	very	bad	but	what	a	great	music
-----	-----	----	------	-----	-----	------	---	-------	-------

Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ $very \rightarrow bad ++$						

Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ <i>very</i> \rightarrow <i>bad</i> ++						
			$\underbrace{\hspace{2.5cm}}$ <i>but</i> will change <i>bad</i>						

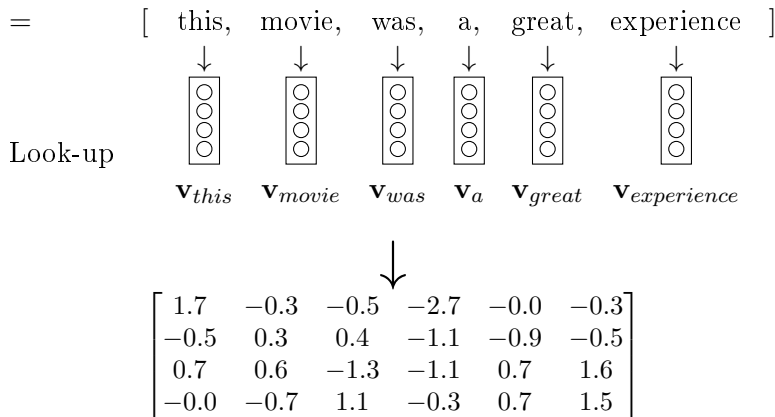
Local contexts

the	end	is	very	bad	but	what	a	great	music
			$very \rightarrow bad++$						
			but will change bad						
		bad is for end not $music$						$great$ is for $music$ not fo end	

Motivations

- Local contextualisation
- Global view of the sentence

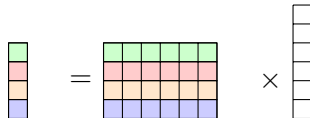
Another view of a sentence



Propose 2 solutions for an improved text classification

...

Matrix and Vector product

$$\mathbf{y} = \mathbf{W} \times \mathbf{x}$$


$$\begin{aligned} y_1 &= \mathbf{W}_{1,:} \times \mathbf{x} \\ y_2 &= \mathbf{W}_{2,:} \times \mathbf{x} \\ \dots &\dots \end{aligned}$$

In terms of dimension:

$$\text{With } \begin{cases} \mathbf{x} & : (L_1 \times 1) \\ \mathbf{W} & : (L_2 \times C_2) \\ \mathbf{y} & : (L_2 \times 1) \end{cases} \Rightarrow (\mathbf{W}\mathbf{x}) : (L_2 \times 1) = (L_2 \times \underbrace{C_2}_{C_2=L_1}) (\cancel{L_1} \times 1)$$

Matrix-matrix product

\mathbf{X} is a matrix of 2 columns, 2 vectors as \mathbf{x} :

$$\mathbf{Y} = \mathbf{W} \times \mathbf{X}$$

$$\begin{bmatrix} \text{green} & \text{red} \\ \text{red} & \text{red} \\ \text{orange} & \text{red} \\ \text{blue} & \text{red} \end{bmatrix} = \begin{bmatrix} \text{green} & \text{green} & \text{green} & \text{green} & \text{green} \\ \text{red} & \text{red} & \text{red} & \text{red} & \text{red} \\ \text{orange} & \text{orange} & \text{orange} & \text{orange} & \text{orange} \\ \text{blue} & \text{blue} & \text{blue} & \text{blue} & \text{blue} \end{bmatrix} \times \begin{bmatrix} \text{white} & \text{red} \\ \text{white} & \text{red} \\ \text{white} & \text{red} \\ \text{white} & \text{red} \\ \text{white} & \text{red} \end{bmatrix}$$

In terms of dimension:

$$\text{With } \begin{cases} \mathbf{X} & : (L_1 \times C_1) \\ \mathbf{W} & : (L_2 \times C_2) \\ \mathbf{y} & : (L_2 \times C_1) \end{cases} \rightarrow (L_2 \times C_2) = (L_2 \times \underbrace{\cancel{C_2}(\cancel{L_1} \times C_1)}_{C_2=L_1})$$

Outline

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Draw attention for classification

Remind CBOW classifier

The classifier output:

$$\text{softmax}(\mathbf{W}^o \mathbf{h}) \text{ (multiclass) or } \sigma(\mathbf{w}^o \mathbf{h}) \text{ (binary)}$$

- What does represent a row of \mathbf{W}^o ?
- The product $\mathbf{W}^o \mathbf{h}$?
- The softmax ?

Draw attention

Is a word vector related to the classification task ?

$$\mathbf{h} = \sum_{i=1}^L \underbrace{\mathbf{x}_i}_{\text{emb. of word } i} \longrightarrow \mathbf{h} = \sum_{i=1}^L \underbrace{\lambda_i}_{\text{???}} \mathbf{x}_i$$

Draw attention for classification (binary task)

$$\mathbf{X}\mathbf{q} = L \left\{ \begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & & & \\ \hline & & & \\ \hline \end{array} \right\} \times \begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array} = \begin{array}{|c|} \hline \\ \hline \\ \hline \\ \hline \\ \hline \end{array} \in \mathbb{R}^L$$
$$(\mathbf{X}\mathbf{q})_i = \mathbf{x}_i^t \mathbf{q} \quad (\text{dot product})$$
$$\mathbf{a} = \text{softmax}(\mathbf{X}\mathbf{q})$$

- $\mathbf{a} = (a_i)$, $\sum_{i=1}^L a_i = 1$ and $0 \leq a_i \leq 1$
- \mathbf{a} : attention vector for the "query" \mathbf{q} and the "keys" \mathbf{X} .
- \mathbf{q} is a vector to be learnt [11, 7]

Attention to weight inputs (binary task)

- $\mathbf{a} = \text{softmax}(\mathbf{X}\mathbf{q})$ is the attention vector

$$\mathbf{h} = \sum_{i=1}^L a_i \mathbf{x}_i = \mathbf{a}^t \mathbf{X}$$

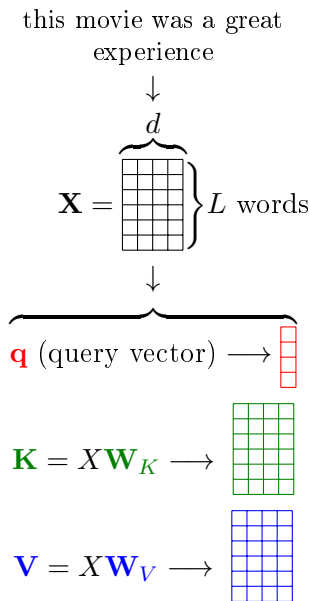
- A new vector, focused on the classification task (\mathbf{q})
- To summarize:

$$\mathbf{h} = \text{softmax}(\mathbf{X}\mathbf{q})^t \mathbf{X} \rightarrow \text{classification}$$

Issues:

- Scale the dot product
- \mathbf{X} is involved everywhere !

Basic attention mechanism for classification (binary task)



$$\mathbf{h} = \text{softmax} \left(\frac{\mathbf{K}\mathbf{q}}{\sqrt{d}} \right)^t \mathbf{V}$$

- \mathbf{X} can be static emb.
- or **contextualized embedding**
- \mathbf{q} is learnt as a target for selection
- $\mathbf{a} = \mathbf{K}\mathbf{q}$: selection in \mathbf{V}

Attention classifier: Going to multiclass

Exercise

- How to modify (parametrize) the model for multiclass classification ?
- Can we add more transformations ?

Outline

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Contextualized word embeddings

Consider the word **driver**:

the audio **driver** is really outdated
the **driver** exceeded the speed limit

The context

The	■	The	■	$\lambda_{2,1}$
audio	■■■	driver	■■■■■	$\lambda_{2,2}$
driver	■■■■■	exceeded	■	$\lambda_{2,3}$
is	■	the	■	$\lambda_{2,4}$
really	■	speed	■■■	$\lambda_{2,5}$
outdated	■■■	limit	■	$\lambda_{2,6}$

Self attention: a first idea

Look at the "correlation" between words (embeddings)

- $\mathbf{X}\mathbf{X}^t$ is a $L \times L$ matrix, stores $(\mathbf{x}_i^t \mathbf{x}_j)$
- The i^{th} row stores the "correlation between" \mathbf{x}_i and all the other words in the sentence
- For $i = 2$, we have the correlations with **driver**
- We can use this correlation as a weight

$$\mathbf{z}_2 = \mathbf{z}_{driver} = \sum_{j=1}^L \underbrace{\lambda_{2,j}}_{\mathbf{x}_2^t \mathbf{x}_j} \mathbf{x}_j$$

More (linear) transformations

Two different Transformations on \mathbf{X}

$$\mathbf{X} \longrightarrow \mathbf{X}\mathbf{W}_Q = \mathbf{Q}$$

$$\mathbf{X} \longrightarrow \mathbf{X}\mathbf{W}_K = \mathbf{K},$$


- with \mathbf{W}_Q and $\mathbf{W}_K \in \mathbb{R}^{d \times d}$
- \mathbf{Q} and \mathbf{K} have the same dimensions as \mathbf{X}

$$\mathbf{A} = \mathbf{Q}\mathbf{K}^t = \underbrace{(\mathbf{Q}_{i,*}\mathbf{K}_{j,*}^t)_{i,j}}_{L \times L} = (\mathbf{q}_i^{\mathbf{k}^j}) = (\lambda_{i,j}),$$

with $\lambda_{i,j}$ the attention on "word" j to generate \mathbf{z}_i

Normalization of attention

Take the row-wise softmax:




$$\sum_j \underbrace{\lambda_{i,j}}_{\text{or } a_{i,j}} = 1 \text{ and } \lambda_{i,j} \geq 0$$

Each row of \mathbf{A} gives a convex combination

Self attention (overview)


Consider the word **driver**:



- $(\lambda_{i,j})$ are the attention coefficients, $\sum_j \lambda_{i,j} = 1$, and
- Reflects the influence of \mathbf{x}_j on \mathbf{x}_i (transformed version)


Transformer : Queries, Keys, Values

the driver exceeded the speed limit

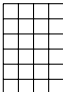


Tranformer : Attention matrix

The distance matrix between Q and K



Scaled Dot-Product Attention

$$\mathbf{Z} = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^t}{\sqrt{d}}\right)\mathbf{V} =$$



Q,K,V and Metric Learning

$$\begin{aligned}\mathbf{Q}\mathbf{K}^t &= \mathbf{X}\mathbf{W}_Q \times (\mathbf{X}\mathbf{W}_K)^t = \mathbf{X}\mathbf{W}_Q \times (\mathbf{W}_K^t \mathbf{X}^t) \\ &= \mathbf{X}\mathbf{M}\mathbf{X}^t\end{aligned}$$

- If \mathbf{M} would be PSD, it is a metric.
- Otherwise, it is a transformed similarity (bilinear similarity)

\mathbf{M} is learnt: a transformer block learns its own similarity.

Multi-head attention (with 2 heads)




Putting all together (with more tricks)

Transformer block

From [10]


- Inputs is \mathbf{X}
- Positional embeddings
- Multihead attention
- Residual connections [6]
- Layer Normalization [2]
- Final filtering



Layer norm


Assume \mathbf{Z} a minibatch of sequences (B, L, D) : $\mathbf{Z} = L \left\{ \begin{array}{c} \text{grid} \\ \vdots \\ \text{grid} \end{array} \right\}$
 d

Batch or Layer norm



[9]

Positional embeddings




- Originally "absolute"
- Can be learnt [5, 1]
- Or relative [8]

(figure generated by the following code


```
https://github.com/jalammar/jalammar.github.io/blob/master/notebooks/  
transformer/transformer\_positional\_encoding\_graph.ipynb)
```


A Transformer layer




Transformer layers can be stacked !

Pre-training as a (Masked) language model



BERT Encoder for text classification



Outline

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Summary

Attention, attention

- This mechanism allows the model to efficiently handle different kind of structure.
- Originally for machine translation, and with BI-GRU [4, 3].

Transformers

- Architecture proposed in [10]
- Nowadays state of the art component

Transformers are everywhere

State of the art encoder

- For text ! (BERT)
- And also for speech, DNA, vision, ...

Also a powerful generator

- For text (GPT, ...)
- Speech, ... sequences

Outline

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

- [1] Rami Al-Rfou et al. *Character-Level Language Modeling with Deeper Self-Attention*. 2018. arXiv: 1808.04444 [cs.CL].
- [2] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. *Layer Normalization*. 2016. arXiv: 1607.06450 [stat.ML].
- [3] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. “Neural Machine Translation by Jointly Learning to Align and Translate”. In: *CoRR* abs/1409.0473 (2014). URL: <http://arxiv.org/abs/1409.0473>.
- [4] Kyunghyun Cho et al. “Learning Phrase Representations using RNN Encoder–Decoder for Statistical Machine Translation”. In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Doha, Qatar: Association for Computational Linguistics, 2014, pp. 1724–1734. URL: <http://www.aclweb.org/anthology/D14-1179>.
- [5] Jonas Gehring et al. “Convolutional Sequence to Sequence Learning”. In: *CoRR* abs/1705.03122 (2017). arXiv: 1705.03122. URL: <http://arxiv.org/abs/1705.03122>.
- [6] Kaiming He et al. “Deep Residual Learning for Image Recognition”. In: *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016, pp. 770–778. arXiv: 1512.03385 [cs.CV].
- [7] Zhouhan Lin et al. “A STRUCTURED SELF-ATTENTIVE SENTENCE EMBEDDING”. In: *International Conference on Learning Representations*. 2017. URL: https://openreview.net/forum?id=BJC_jUqxe.

- [8] Peter Shaw, Jakob Uszkoreit, and Ashish Vaswani. *Self-Attention with Relative Position Representations*. 2018. arXiv: 1803.02155 [cs.CL].
- [9] Sheng Shen et al. “PowerNorm: Rethinking Batch Normalization in Transformers”. In: *Proceedings of the 37th International Conference on Machine Learning*. Ed. by Hal Daumé III and Aarti Singh. Vol. 119. Proceedings of Machine Learning Research. PMLR, 2020, pp. 8741–8751. URL: <https://proceedings.mlr.press/v119/shen20e.html>.
- [10] Ashish Vaswani et al. “Attention is All you Need”. In: *Advances in Neural Information Processing Systems 30*. Ed. by I. Guyon et al. Curran Associates, Inc., 2017, pp. 6000–6010. URL: <http://papers.nips.cc/paper/7181-attention-is-all-you-need.pdf>.
- [11] Zichao Yang et al. “Hierarchical Attention Networks for Document Classification”. In: *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)06*. 2016.