

# 6 – Scaling LLMs

IASD Apprentissage – LLMs course

Florian Le Bronnec

December 3, 2024

# Table of Contents

- ① Exam
- ② Last time
- ③ Big models
  - Scaling the data
  - Scaling the models
- ④ Scaling the context
  - Decreasing the cost of attention
- ⑤ Chat Models

# Table of Contents

- ① Exam
- ② Last time
- ③ Big models
  - Scaling the data
  - Scaling the models
- ④ Scaling the context
  - Decreasing the cost of attention
- ⑤ Chat Models

# Exam

## Information

- Date: 10/12
- 2h.
- No computer, no internet.
- 2 pages (recto-verso) of handwritten notes.
- Exam will feature questions about an article AND the course.
- The article will be made available on the 03/12 at 12h.
- No article during the exam.

## Advice

- No need to learn by heart.
- If a question is about a specific point of the article, it will be reminded in the exam subject.

# Mock exam on Attention is all you need

## Question 1: Encoder Input in Translation

**What is the input to the encoder in translation?**

## Question 1: Encoder Input in Translation

**What is the input to the encoder in translation?**

**Answer:** The input to the encoder is a sequence of symbol representations (e.g., embeddings) of the source sentence. These embeddings are combined with positional encodings to inject information about the order of tokens in the sequence.

## Question 2: Decoder Input in Translation

**During training, what is the input of the decoder? What are its labels?**



## Question 2: Decoder Input in Translation

**During training, what is the input of the decoder? What are its labels?**

**Answer:** The input to the decoder is the target sequence's embeddings, combined with positional encodings. The labels are the target sequence shifted by one position to the right.

## Question 3: Step-by-Step Operations in the Decoder

**Describe step-by-step the operations performed in the decoder.**

## Question 3: Step-by-Step Operations in the Decoder

**Describe step-by-step the operations performed in the decoder.**

**Answer:**

- Input Preparation (shifted target sequence + word embeddings + positional encodings).
- Masked Multi-Head Self-Attention
- Encoder-Decoder Multi-Head Attention
- Feed-Forward Network
- Residual Connections and Layer Normalization
- Output Projection

## Question 4: ROUGE and BLEU Metrics and Its Limitations

**Without giving a precise algorithm, how do the BLEU and ROUGE metric work? What are their limitations?**

## Question 4: ROUGE and BLEU Metrics and Its Limitations

**Without giving a precise algorithm, how do the BLEU and ROUGE metric work? What are their limitations?**

**Answer:**

- **How it works:**
  - BLEU and ROUGE measure the overlap of n-grams (subsequences of tokens) between a generated sequence and one or more reference sequences.
- **Limitations:**
  - BLEU and ROUGE do not consider semantic meaning, focusing only on surface-level token matches.

## Question 5: Sampling Decoding Algorithm

**Describe in detail the sampling decoding algorithm.**

## Question 5: Sampling Decoding Algorithm

Describe in detail the sampling decoding algorithm.

Answer:

- Initialize the generated sequence with a start token <BOS>.
- Repeat until the end token <EOS> is generated or a maximum length is reached:
  - **Step 1: Input Preparation.** Feed the current sequence into the model to obtain the probability distribution  $P(y_t \mid \text{previous tokens})$  over the vocabulary.
  - **Step 2: Sampling.** Use the probability distribution to sample the next token  $y_t$ . This can involve:
    - Random multinomial sampling from  $P(y_t)$ .
  - **Step 3: Append Token.** Add the sampled token to the sequence.
- Return the generated sequence.

# Table of Contents

- ① Exam
- ② **Last time**
- ③ Big models
  - Scaling the data
  - Scaling the models
- ④ Scaling the context
  - Decreasing the cost of attention
- ⑤ Chat Models



# Last time – Finetuning

## Finetuning of encoder models

Encoder (BERT-like models) are tuned for specific tasks.

- **Classification:** add a classification head on top of the model, on the [CLS] token.
- **Token classification (or tagging):** add a classification head on top of the model, on each token.
- **Sequence pairs classification:** add a classification head on top of the model, on the [SEP] token.

**Tasks:** NER, POS tagging, sentiment analysis, etc.

# Last time – Finetuning

## Finetuning of decoder models

Decoder (GPT-like models) are tuned for specific tasks.

- Models are still trained with a **next token prediction** objective.
- Concatenate the source with the target, and train the model to predict the target.

**Tasks:** translation, summarization, question answering, etc.

# Table of Contents

- ① Exam
- ② Last time
- ③ Big models
  - Scaling the data
  - Scaling the models
- ④ Scaling the context
  - Decreasing the cost of attention
- ⑤ Chat Models

# What's under the hood?

## Impressive pretrained models

- BERT-like models achieve **better performances than humans** on the SuperGLUE dataset [6].
- BART-like models can **summarize texts, answer questions**, etc. [2].
- GPT3-3.5-4 are few-shot learners, they can answer questions based on their **general knowledge** [4].
- ChatGPT.

# What's under the hood?

## Impressive pretrained models

- BERT-like models achieve **better performances than humans** on the SuperGLUE dataset [6].
- BART-like models can **summarize texts, answer questions**, etc. [2].
- GPT3-3.5-4 are few-shot learners, they can answer questions based on their **general knowledge** [4].
- ChatGPT.

## How?

- Data.
- Scaling the models.

# What's under the hood?

## Impressive pretrained models

- BERT-like models achieve **better performances than humans** on the SuperGLUE dataset [6].
- BART-like models can **summarize texts, answer questions**, etc. [2].
- GPT3-3.5-4 are few-shot learners, they can answer questions based on their **general knowledge** [4].
- ChatGPT.

## How?

- Data.
- Scaling the models.

⇒ What does it represent in practice?

# Data

Model	Params	Context	Batch	Steps
BERT [1]	355M	512	256 <sup>1</sup>	1M
BART [2]	406M	1024	8000	500K
LLama2 [11]	7-13-70B	4096	4000 <sup>1</sup>	500K

Table 1: Comparison of NLP models. Steps are the number of training steps.

<sup>1</sup>Batch size were indicated in terms of number of tokens, I approximately converted it to number of documents.

# Data

Model	Params	Context	Batch	Steps
BERT [1]	355M	512	256 <sup>1</sup>	1M
BART [2]	406M	1024	8000	500K
LLama2 [11]	7-13-70B	4096	4000 <sup>1</sup>	500K

Table 1: Comparison of NLP models. Steps are the number of training steps.

Demo!

---

<sup>1</sup>Batch size were indicated in terms of number of tokens, I approximately converted it to number of documents.



# Pretraining in practice

Biggest NVIDIA GPUs are  $\sim 8\times$  bigger than Colab's GPUs.

Even with bigger GPUs, processing batches of **8000 documents** is an issue.  
How can we do that in practice?

# Pretraining in practice

Biggest NVIDIA GPUs are  $\sim 8\times$  bigger than Colab's GPUs.

Even with bigger GPUs, processing batches of **8000 documents** is an issue.  
How can we do that in practice?

$\implies$  Let's review some methods to perform these heavy trainings.

# Data parallelism




Figure 1: Data parallelism.

# Gradient accumulation

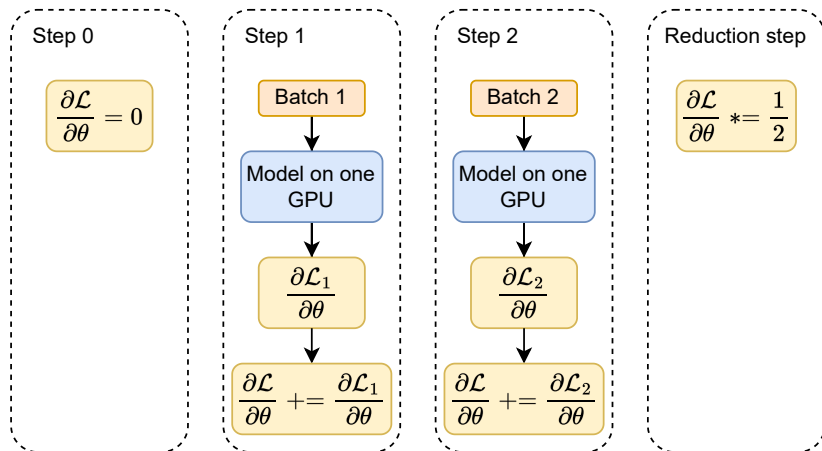


Figure 2: Gradient accumulation.

# Summary

**Data parallelism** and **gradient accumulation** are two very common methods in NLP (and deep learning) to increase the effective size of the batch.

## Data Parallelism

Per Device Batch Size	Number of Devices	Effective Batch Size
32	1	32
32	2	64
32	4	128

## Gradient accumulation

Batch Size	Accumulation Steps	Effective Batch Size
32	1	32
32	2	64
32	4	128

# Comparison

Aspect	Data Parallelism	Gradient Accumulation
Parallelism	Yes, across GPUs	No, sequential computations
Several GPUs	Yes	Can work on 1 GPU
Time	$1 \times (\text{forward} + \text{backward})$ + communication across GPUs + parameters update	$N \times (\text{forward} + \text{backward})$ + parameters update

**Table 2:** Comparison of Data Parallelism and Gradient Accumulation.  $N$  is the number of accumulation steps.

# Comparison

Aspect	Data Parallelism	Gradient Accumulation
Parallelism	Yes, across GPUs	No, sequential computations
Several GPUs	Yes	Can work on 1 GPU
Time	$1 \times (\text{forward} + \text{backward})$ + communication across GPUs + parameters update	$N \times (\text{forward} + \text{backward})$ + parameters update

**Table 2:** Comparison of Data Parallelism and Gradient Accumulation.  $N$  is the number of accumulation steps.

⇒ Data parallelism is useful when you have **several GPUs** and want to **speed up** the training.

⇒ Gradient accumulation is useful when you have **limited resources**.

Both can of course be combined (and are combined in practice).

# In practice?

Everything is quite easy with PyTorch.

Demo!



# Scaling the number of parameters

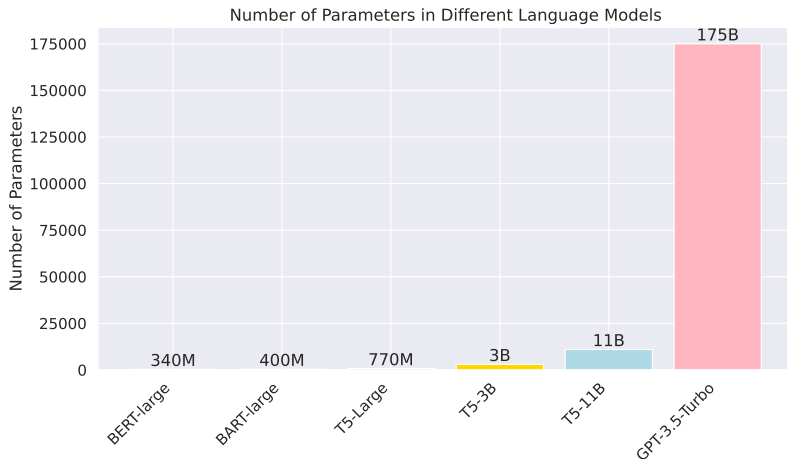


Figure 3: Different models and their scales.

# What does it imply?

Demo on GPT2 size.

# What does it imply?

Demo on GPT2 size.

GPT2-Large on a GPU is  $\sim 3.5\text{GiB}$ . GPT-3 is  $\sim \times 250$  bigger than GPT2-Large.

Biggest GPUs are 80GiB.

$\implies$  How does it fit?

# What does it imply?

Demo on GPT2 size.

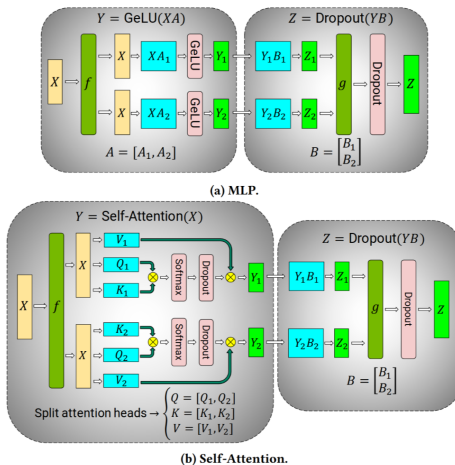
GPT2-Large on a GPU is  $\sim 3.5\text{GiB}$ . GPT-3 is  $\sim \times 250$  bigger than GPT2-Large.

Biggest GPUs are 80GiB.

$\implies$  How does it fit?

We're going to speak about **training** and **finetuning** / **inference**.

# Model parallelism



1

Figure 4: Idea: split the computations across several GPUs. Tensor parallelism, figure from [5].

# Pretraining

For pretraining, there are some heavy hardware optimizations, like the ones presented in [5].

The idea is to split independant and costly operations across devices, then aggregate the final result.

Even if there are solutions, as a rule of thumb remind that **model parallelism is not easy**.

# And for end-users?

Most people are not interested in what we presented in previous slides.

But still, we might want to use these models for at least:

- running inference as is,
- finetuning on a specific task.

⇒ Let's describe some solutions for practitioners.

## Some solutions: Quantization

GPUs standard precision is `float32`. A solution is to **reduce this precision**.

Data Type	Bit Width	Hardware Capability	Use for training
<code>float32</code>	32	General-purpose CPUs/GPUs	Yes
<code>float16</code>	16	GPUs with FP16 support	Yes
<code>bfloat16</code>	16	NVIDIA Ampere GPUs, TPUs	Yes
<code>int8</code>	8	CPUs, GPUs	No

**Table 3:** Non-exhaustive list of mixed-precision data types and hardware support.



## Some solutions: Quantization

GPUs standard precision is `float32`. A solution is to **reduce this precision**.

Data Type	Bit Width	Hardware Capability	Use for training
<code>float32</code>	32	General-purpose CPUs/GPUs	Yes
<code>float16</code>	16	GPUs with FP16 support	Yes
<code>bfloat16</code>	16	NVIDIA Ampere GPUs, TPUs	Yes
<code>int8</code>	8	CPUs, GPUs	No

**Table 3:** Non-exhaustive list of mixed-precision data types and hardware support.

⇒ In practice the models maintain their performances (empirical statement).

### Be careful when using quantization

- Make sure the operations you use are well supported.
- Read the documentation, especially for training.
- There can be some instabilities (ex: T5 does not work with `float16`).

# Reduce the cost of autodiff

Component	Memory Cost	Inference
Model Parameters	$O(P)$	Yes
Activations	$O(B \times L)$	Yes / No
Gradients	$O(P)$	No
Optimizer States	$O(P)$	No

Table 4: Memory cost during training and inference.

Huge dependency on  $P$ , the number of parameters to update.

# Reduce the cost of autodiff

Component	Memory Cost	Inference
Model Parameters	$O(P)$	Yes
Activations	$O(B \times L)$	Yes / No
Gradients	$O(P)$	No
Optimizer States	$O(P)$	No

Table 4: Memory cost during training and inference.

Huge dependency on  $P$ , the number of parameters to update.

⇒ Reduce the number of parameters to update!

# Parameter efficient finetuning

Let  $\theta \in \mathbb{R}^P$  be the parameters of a base model: BERT, Llama, etc.

The goal of **parameter efficient finetuning (PEFT)** is to select a subset of parameters  $\theta' \in \mathbb{R}^{P'}$  with  $P' \ll P$  such that the model's performances are maintained.

## Example

Train only the last layer of the model.

Break!

# Table of Contents

- ① Exam
- ② Last time
- ③ Big models
  - Scaling the data
  - Scaling the models
- ④ Scaling the context
  - Decreasing the cost of attention
- ⑤ Chat Models

# Which length of texts can we process?

**How many tokens** can fit in the models?

There are 3 types of limitations:

- **architectural**, because of position embeddings of fixed size,
- **training setup**, the maximal length seen during training, independently of the position embeddings,
- **computational cost**, because the costs scale with the context length.

# Attention is $\mathcal{O}(L^2)$

Let  $L$  be the input length.

$$QK^T \in \mathbb{R}^{L \times L}.$$

Demo!



# Attention is $\mathcal{O}(L^2)$

Let  $L$  be the input length.

$$QK^T \in \mathbb{R}^{L \times L}.$$

Demo!

$\Rightarrow$  **Quadratic** cost w.r.t  $L$ .

How can we avoid that?

# Scaling the context

Simple idea: **reduce the size of the attention matrix.**




Figure 5: Reducing the size of the attention matrix.

We can therefore reach a  $\mathcal{O}(Lw)$  memory cost.

# Sparse attention




Figure 6: Sparse attention patterns, figure from [7].

Whole literature on sparse transformers:

- LongT5 [9],
- BigBird [7],
- Longformer [3] (extends pretrained models),
- etc.

# Other solutions

## Hardware improvements

- FlashAttention [8] proposed recently a new CUDA kernel optimized for computing attention on recent NVIDIA GPUs.

# Other solutions

## Hardware improvements

- FlashAttention [8] proposed recently a new CUDA kernel optimized for computing attention on recent NVIDIA GPUs.

## Stick to the model's base length

- 4k tokens is already a lot.
- Models have **not been pretrained on such lengths**.
- LLMs **do not use their full context** [10].
- Use **retrieval techniques**.
- Process the documents by **chunks** (e.g. summarize chapter by chapter).

# What about other architectures?

New architectures are being proposed

State-space models, MAMBA, etc.

# Table of Contents

- ① Exam
- ② Last time
- ③ Big models
  - Scaling the data
  - Scaling the models
- ④ Scaling the context
  - Decreasing the cost of attention
- ⑤ Chat Models

# Introduction to Training Chat Models

- **Goal:** Develop models capable of understanding and responding to user instructions in a conversational manner.
- **Challenge:** Achieving natural, helpful, and aligned responses requires more than basic fine-tuning.
- **Two-Phase Training Approach:**
  - **Phase 1 – Instruction-Tuning:** Fine-tune the model on instruction-based datasets to improve task understanding.
  - **Phase 2 – Reinforcement Learning from Human Feedback (RLHF):** Use human feedback to further align the model's responses with user preferences.
- This combination makes chat models able of going beyond basic instruction-following, incorporating nuanced human feedback.



# Instruction-Tuning for Chat Models

- **Objective:** Adapt generative models to follow user instructions accurately.
- **Method:** Fine-tune models on datasets where each example includes an instruction and a correct response.
- Examples of instruction-tuning datasets: FLAN, Super-NaturalInstructions.

## Example of Supervised Instruction-Tuning

- **Task:** Fine-tune the model to respond accurately to user instructions.

Instruction	Model Input	Output Prediction
Instruction: "List three benefits of regular exercise." Expected Output: "Improves cardiovascular health, boosts mental well-being, and strengthens muscles."	[BOS] List three benefits of regular exercise.[EOS]	Improves cardiovascular health, boosts mental well-being, and strengthens muscles.

- **Objective:** Model learns to generate responses aligned with the instruction prompt.

# Limitations of Instruction-Tuning

- **Instruction-Tuning Provides a Solid Start:**
  - Models are trained to follow directions and respond to a wide range of prompts.
  - Instruction datasets enable models to generalize across many basic tasks.

# Limitations of Instruction-Tuning

- **Instruction-Tuning Provides a Solid Start:**
  - Models are trained to follow directions and respond to a wide range of prompts.
  - Instruction datasets enable models to generalize across many basic tasks.
- **Key Limitations of Instruction-Tuning:**
  - **Models are trained using only next token prediction objective:** Models will learn the dataset distribution. This does not encompass directly any notion of "quality", like helpfulness, factuality, toxicity, etc.
  - **Datasets are not perfect:** After learning, models might reflect some imperfections of the training data, like generic tone or writing style that might not be appropriate in all contexts.
  - **No filtering on sensitive domains:** Instruction datasets are often filtered, and often contain generic or neutral content. Therefore, models should be further trained to handle sensitive domains, like avoiding harmful content.

# Reinforcement Learning from Human Feedback (RLHF)

- **Solution:**

# Reinforcement Learning from Human Feedback (RLHF)

- **Solution:** Improve response quality and relevance through direct human feedback, guiding models to align closer with human expectations.
- **Goal of RLHF:**
  - Use human feedback to enhance the model's ability to produce "human-like" responses.

# Reinforcement Learning from Human Feedback (RLHF)

- **Solution:** Improve response quality and relevance through direct human feedback, guiding models to align closer with human expectations.
- **Goal of RLHF:**
  - Use human feedback to enhance the model's ability to produce "human-like" responses.

# Why RLHF Matters

- **Why RLHF Matters:**
  - Aligns model responses more closely with human-like preferences and expectations.
  - Provides a method for refining responses beyond the limitations of instruction-tuning.
- **Examples of RLHF Datasets:** Anthropic's HH-RLHF, OpenAI's feedback datasets.



# Correcting Model Behavior with Human Feedback

**Objective:** Improve model responses by correcting cases where generated outputs are suboptimal or misaligned with user expectations.

## Ideal Approach:

- 1 Generate a variety of responses to a prompt.
- 2 Make a human rate each response, on a scale from 0 to 5.
- 3 Use a reinforcement learning (RL) algorithm to adjust the model's behavior, encouraging preferred responses and discouraging low-quality ones.

Such RL algorithms can be REINFORCE, PPO, etc.

These methods might help correct and refine model behavior, aligning responses more closely with human expectations.

# Why Direct Human Ratings are Impractical

**Challenge with Human Ratings:** While human ratings provide valuable feedback, rating every response generated by the model is unrealistic.

## Limitations of Relying on Human Ratings:

- **High Cost and Time Demand:** Scaling human feedback for all responses is resource-intensive.
- **Slow Iteration:** Relying on human feedback would slows down model training.

To make reinforcement learning feasible at scale, we need a way to approximate these ratings efficiently.

# Approximating Human Feedback with a Value Model

Instead of using human, we will train a model to approximate human preferences. **Step 1: Collect human feedback:**

- For a given prompt, we generate 2 responses.
- Human annotators are asked to rank the responses based on quality (a preferred response and a less-preferred response).
- We train the value model (typically a value head on top of the finetuned LM) to predict whether a response is a good one or not (independently of each other), using standard binary classification loss.

After training, we have a value model  $r$  which is able to give a "score" to a response,  $r(y)$ , which corresponds to the probability of the response being the preferred one.

# Using a Value Model with REINFORCE

**Goal:** Use the value model to approximate rewards for generated responses, guiding the main model's updates.

## REINFORCE Algorithm with Value Model:

- For a prompt  $x$ , the model generates a response  $y$ .
- The value model assigns a score  $r(y)$  to the response, approximating human preference.
- Update the main model's parameters  $\theta$  to maximize expected reward  $\mathbb{E}[r(y)]$  using the gradient:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{y \sim \pi_{\theta}(y|x)} [r(y) \nabla_{\theta} \log \pi_{\theta}(y|x)]$$

# RLHF – Summary

- Learn a **value model** to approximate human feedback scores.
- Use **REINFORCE** or **PPO** to update the main model based on the value model's scores.
  - Generate responses to intrusions.
  - Compute the score using the value model.
  - Update the main model using REINFORCE or PPO.

# Motivation for Direct Preference Optimization (DPO)

## Limitations of RL with Value Models:

- REINFORCE or PPO require a value model to approximate human feedback scores.
- Training with RL involves balancing multiple components, such as reward scaling and policy stability, making it complex and computationally intensive.

## Can we get rid of RL?:

- Yes! Use directly the important piece of information: preference dataset.
- This is the idea behind Direct Preference Optimization (DPO).

# Transition from RL to DPO

In RL (e.g., REINFORCE), we maximize the expected reward:

$$J(\theta) = \mathbb{E}_{y \sim \pi_{\theta}(y|x)} [r(y)]$$

In DPO, rather than optimizing for reward scores, we optimize preferences directly. For a prompt  $x$  and two responses  $y_{\text{preferred}}$  and  $y_{\text{less\_preferred}}$ , DPO adjusts the model to make  $y_{\text{preferred}}$  more likely.

**DPO Loss:**

$$\mathcal{L}_{\text{DPO}} = -\log \sigma \left( \log \frac{\pi_{\theta}(y_{\text{preferred}}|x)}{\pi_{\theta}(y_{\text{less\_preferred}}|x)} \right)$$

where  $\sigma$  is the sigmoid function.

# Complete RL and DPO objective

## Regularized RL Loss with Reference Model:

- In reinforcement learning, we often use a regularization term to keep the fine-tuned model  $\pi_\theta$  close to a base reference model  $\pi_{\text{ref}}$ , which helps maintain stability and prevents excessive deviation.

$$\mathcal{L}_{\text{RL}} = -\mathbb{E}_{y \sim \pi_\theta(y|x)} [r(y)] - \lambda \text{KL}(\pi_\theta \| \pi_{\text{ref}})$$

## Regularized DPO Loss with Reference Model:

- In DPO, this regularization translate to using a base model  $\pi_{\text{ref}}$  to stabilize training (full proof in original DPO paper).

$$\mathcal{L}_{\text{DPO}} = -\log \sigma \left( \log \frac{\pi_\theta(y_{\text{preferred}}|x)/\pi_{\text{ref}}(y_{\text{preferred}}|x)}{\pi_\theta(y_{\text{less\_preferred}}|x)/\pi_{\text{ref}}(y_{\text{less\_preferred}}|x)} \right)$$

where  $\pi_{\text{ref}}$  keeps  $\pi_\theta$  anchored to the reference model's distribution.



# Contrastive Learning in DPO – Introduction

## Contrastive Learning Overview:

- In contrastive learning, the goal is to make similar samples closer while pushing dissimilar samples apart in the model's representation space.

## Applying Contrastive Learning to Preferences:

- In DPO, we have a prompt  $x$  with two responses:
  - $y_{\text{preferred}}$ : the response preferred by human feedback.
  - $y_{\text{less\_preferred}}$ : a lower-rated alternative.
- The objective is to make  $y_{\text{preferred}}$  more likely than  $y_{\text{less\_preferred}}$ , similar to bringing closer and pushing apart in contrastive learning.

# DPO as a Contrastive Objective

## DPO Loss as a Contrastive Loss:

- The DPO loss function:

$$\mathcal{L}_{\text{DPO}} = -\log \sigma \left( \log \frac{\pi_{\theta}(y_{\text{preferred}}|x)}{\pi_{\theta}(y_{\text{less\_preferred}}|x)} \right)$$

optimizes the model to assign higher probability to preferred responses.

- This can be interpreted as a contrastive objective:
  - Maximizing  $\pi_{\theta}(y_{\text{preferred}}|x)$  pulls preferred responses “closer” (higher probability).
  - Minimizing  $\pi_{\theta}(y_{\text{less\_preferred}}|x)$  pushes less-preferred responses “farther” (lower probability).

## Conclusion: DPO vs. RL for Human Alignment

- **Simplicity:** DPO directly optimizes for preferences, avoiding the complexity of reward modeling in RL.
- **Efficiency:** DPO's contrastive objective is computationally lighter, allowing faster, more stable training.

## Conclusion: DPO vs. RL for Human Alignment

- **Simplicity:** DPO directly optimizes for preferences, avoiding the complexity of reward modeling in RL.
- **Efficiency:** DPO's contrastive objective is computationally lighter, allowing faster, more stable training.

### Performance Gains from Human Alignment:

- DPO and other human-alignment methods have significantly improved model quality, with benchmarks showing up to 20-30% gains in user satisfaction and response relevance.
- DPO now sets the standard for aligning large language models with human feedback, achieving higher consistency and responsiveness.

# Conclusion

Thank you!

# References I

- [1] Jacob Devlin et al. *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*. 2019. [arXiv: 1810.04805 \[cs.CL\]](#).
- [2] Mike Lewis et al. *BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension*. 2019. [arXiv: 1910.13461 \[cs.CL\]](#).
- [3] Iz Beltagy, Matthew E. Peters, and Arman Cohan. *Longformer: The Long-Document Transformer*. 2020. [arXiv: 2004.05150 \[cs.CL\]](#).
- [4] Tom B. Brown et al. *Language Models are Few-Shot Learners*. 2020. [arXiv: 2005.14165 \[cs.CL\]](#).
- [5] Mohammad Shoeybi et al. *Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism*. 2020. [arXiv: 1909.08053 \[cs.CL\]](#).
- [6] Pengcheng He et al. *DeBERTa: Decoding-enhanced BERT with Disentangled Attention*. 2021. [arXiv: 2006.03654 \[cs.CL\]](#).

## References II

- [7] Manzil Zaheer et al. *Big Bird: Transformers for Longer Sequences*. 2021. arXiv: 2007.14062 [cs.LG].
- [8] Tri Dao et al. *FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness*. 2022. arXiv: 2205.14135 [cs.LG].
- [9] Mandy Guo et al. *LongT5: Efficient Text-To-Text Transformer for Long Sequences*. 2022. arXiv: 2112.07916 [cs.CL].
- [10] Mathieu Ravaut et al. *On Position Bias in Summarization with Large Language Models*. 2023. arXiv: 2310.10570 [cs.CL].
- [11] Hugo Touvron et al. *Llama 2: Open Foundation and Fine-Tuned Chat Models*. 2023. arXiv: 2307.09288 [cs.CL].